

# Final Project - Predicting survival of people with Heart Failure

## Libraries

```
In [1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn import linear_model
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import metrics

from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC

from scipy.cluster.hierarchy import linkage, fcluster
from sklearn.cluster import KMeans, DBSCAN
```

## Data Import

```
In [2]: data = pd.read_csv('heart_failure_clinical_records_dataset.csv')
data.head()
```

Out[2]:

	age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction	high_blood_pressure	platelets	serum_creatinine	serum_sodium	sex
0	75.0	0	582	0	20	1	265000.00	1.9	130	1
1	55.0	0	7861	0	38	0	263358.03	1.1	136	1
2	65.0	0	146	0	20	0	162000.00	1.3	129	1
3	50.0	1	111	0	20	0	210000.00	1.9	137	1
4	65.0	1	160	1	20	0	327000.00	2.7	116	0



```
In [3]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 299 entries, 0 to 298
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
---  -
0   age                  299 non-null    float64
1   anaemia              299 non-null    int64
2   creatinine_phosphokinase  299 non-null    int64
3   diabetes             299 non-null    int64
4   ejection_fraction    299 non-null    int64
5   high_blood_pressure    299 non-null    int64
6   platelets            299 non-null    float64
7   serum_creatinine      299 non-null    float64
8   serum_sodium          299 non-null    int64
9   sex                  299 non-null    int64
10  smoking              299 non-null    int64
11  time                  299 non-null    int64
12  DEATH_EVENT           299 non-null    int64
dtypes: float64(3), int64(10)
memory usage: 30.5 KB
```

Number of variables: 13

Types of variables: int64, float64

Redundant Variables: None

```
In [4]: data.dropna(0)
```

Out[4]:

	age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction	high_blood_pressure	platelets	serum_creatinine	serum_sodium	se
0	75.0	0	582	0	20	1	265000.00	1.9	130	:
1	55.0	0	7861	0	38	0	263358.03	1.1	136	:
2	65.0	0	146	0	20	0	162000.00	1.3	129	:
3	50.0	1	111	0	20	0	210000.00	1.9	137	:
4	65.0	1	160	1	20	0	327000.00	2.7	116	(
...	...	...	...	...	...	...	...	...	...	..
294	62.0	0	61	1	38	1	155000.00	1.1	143	:
295	55.0	0	1820	0	38	0	270000.00	1.2	139	(
296	45.0	0	2060	1	60	0	742000.00	0.8	138	(
297	45.0	0	2413	0	38	0	140000.00	1.4	140	:
298	50.0	0	196	0	45	0	395000.00	1.6	136	:

299 rows × 13 columns



```
In [5]: data.shape
```

Out[5]: (299, 13)

*There are no missing data, noise, outliers*


```
In [6]: # statistic description for heart failure data
data.describe().transpose().round(2)
```

Out[6]:


	count	mean	std	min	25%	50%	75%	max
<b>age</b>	299.0	60.83	11.89	40.0	51.0	60.0	70.0	95.0
<b>anaemia</b>	299.0	0.43	0.50	0.0	0.0	0.0	1.0	1.0
<b>creatinine_phosphokinase</b>	299.0	581.84	970.29	23.0	116.5	250.0	582.0	7861.0
<b>diabetes</b>	299.0	0.42	0.49	0.0	0.0	0.0	1.0	1.0
<b>ejection_fraction</b>	299.0	38.08	11.83	14.0	30.0	38.0	45.0	80.0
<b>high_blood_pressure</b>	299.0	0.35	0.48	0.0	0.0	0.0	1.0	1.0
<b>platelets</b>	299.0	263358.03	97804.24	25100.0	212500.0	262000.0	303500.0	850000.0
<b>serum_creatinine</b>	299.0	1.39	1.03	0.5	0.9	1.1	1.4	9.4
<b>serum_sodium</b>	299.0	136.63	4.41	113.0	134.0	137.0	140.0	148.0
<b>sex</b>	299.0	0.65	0.48	0.0	0.0	1.0	1.0	1.0
<b>smoking</b>	299.0	0.32	0.47	0.0	0.0	0.0	1.0	1.0
<b>time</b>	299.0	130.26	77.61	4.0	73.0	115.0	203.0	285.0
<b>DEATH_EVENT</b>	299.0	0.32	0.47	0.0	0.0	0.0	1.0	1.0

```
In [7]: ax = sns.distplot(data[data['DEATH_EVENT'] == 1][['age']])  
ax.figure.set_size_inches(20, 5)  
ax.set(title = 'Age vs Death', xlabel = 'age', ylabel = 'deaths')
```

```
Out[7]: [Text(0, 0.5, 'deaths'), Text(0.5, 0, 'age'), Text(0.5, 1.0, 'Age vs Death')]
```

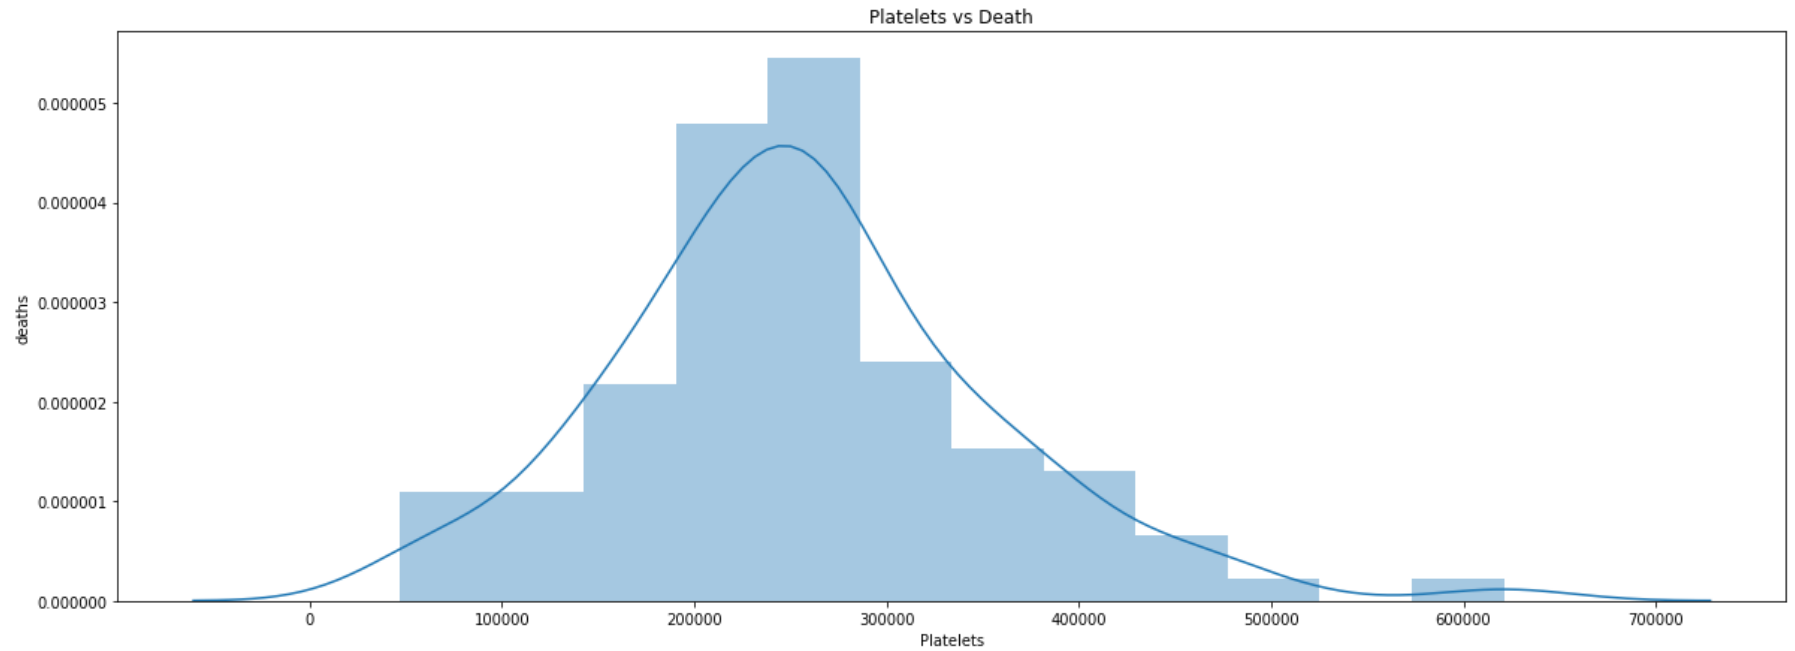


```
In [8]: fig, axs = plt.subplots(1)
fig.set_size_inches(15, 5)
axs = sns.distplot(data[data['DEATH_EVENT'] == 1][['creatinine_phosphokinase']])
```




```
In [9]: ax = sns.distplot(data[data['DEATH_EVENT'] == 1][['platelets']])
ax.figure.set_size_inches(20, 7)
ax.set(title = 'Platelets vs Death', xlabel = 'Platelets', ylabel = 'deaths')
```

```
Out[9]: [Text(0, 0.5, 'deaths'),
Text(0.5, 0, 'Platelets'),
Text(0.5, 1.0, 'Platelets vs Death')]
```



```
In [10]: ax = sns.distplot(data[data['DEATH_EVENT'] == 1][['serum_creatinine']])
ax.figure.set_size_inches(20, 5)
ax.set(title = 'Serum Creatinine vs Death', xlabel = 'Serum Creatinine', ylabel = 'deaths')
```

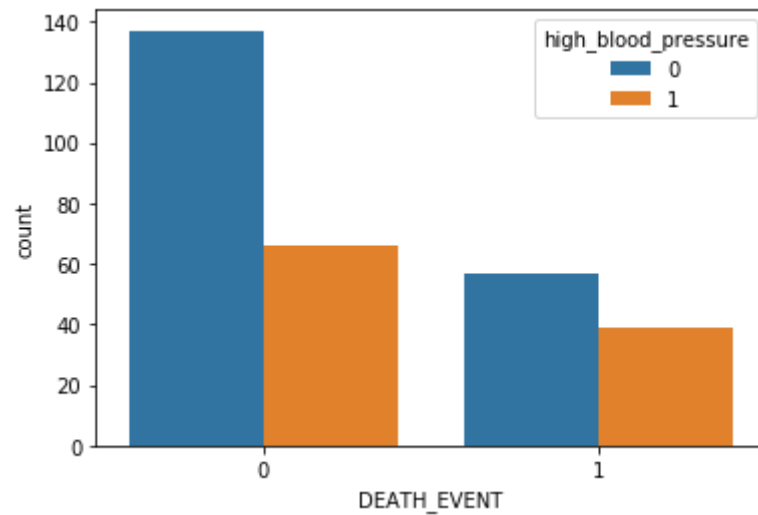
```
Out[10]: [Text(0, 0.5, 'deaths'),
Text(0.5, 0, 'Serum Creatinine'),
Text(0.5, 1.0, 'Serum Creatinine vs Death')]
```






```
In [11]: sns.countplot(x = 'DEATH_EVENT', hue = 'high_blood_pressure', data = data)
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x7faa66ed0890>
```




```
In [12]: sns.countplot(x = 'DEATH_EVENT', hue = 'diabetes', data = data)
```

```
Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x7faa66e3d450>
```




```
In [13]: sns.countplot(x = 'anaemia', hue = 'DEATH_EVENT', data = data)
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7faa66dbd050>
```



```
In [14]: sns.countplot(x = 'sex', hue = 'DEATH_EVENT', data = data)
```

```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x7faa66d25150>
```



# Data Preparation

```
In [15]: X_train, X_val, y_train, y_val = train_test_split(data.iloc[:, :-1], data.iloc[:, [-1]], test_size = 0.25,
random_state = 0)
# X_train, X_val, y_train, y_val = train_test_split(data.iloc[:, :12], data.iloc[:, 12:], test_size = 0.2
5, random_state = 0)
```

```
In [16]: scaler = StandardScaler()
scaler.fit(X_train)
x_train_scaled = scaler.transform(X_train)
x_val_scaled = scaler.transform(X_val)
# x_val = X_val.select_dtypes(include=[np.int64, np.float64])
x_train_scaled_df = pd.DataFrame(x_train_scaled, index = X_train.index, columns=X_train.columns)
x_val_scaled_df = pd.DataFrame(x_val_scaled, index = X_val.index, columns=X_val.columns)
```

## Classification

### 1: Decision Tree


#### 1.1: Using all predictor variables

```
In [17]: classifier = DecisionTreeClassifier(criterion='entropy', random_state=0)
classifier.fit(x_train_scaled, y_train)
```

```
Out[17]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='entropy',
max_depth=None, max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort='deprecated',
random_state=0, splitter='best')
```

```
In [18]: y_pred = classifier.predict(x_val_scaled)
```

```
In [19]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [20]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy: 0.8133333333333334
Error: 0.18666666666666665
Precision: 0.7241379310344828
Recall: 0.7777777777777778
F1 score: 0.75
```

## 1.2: Using Age and Gender as predictor variables


```
In [21]: selected_predictor_variables = ['age',  
                                         'sex']
```

```
In [22]: classifier = DecisionTreeClassifier(criterion='entropy', random_state=0)  
classifier.fit(x_train_scaled_df[selected_predictor_variables], y_train)
```

```
Out[22]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='entropy',  
                                max_depth=None, max_features=None, max_leaf_nodes=None,  
                                min_impurity_decrease=0.0, min_impurity_split=None,  
                                min_samples_leaf=1, min_samples_split=2,  
                                min_weight_fraction_leaf=0.0, presort='deprecated',  
                                random_state=0, splitter='best')
```

```
In [23]: y_pred = classifier.predict(x_val_scaled_df[selected_predictor_variables])
```

```
In [24]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [25]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy: 0.6533333333333333
Error: 0.3466666666666667
Precision: 0.5333333333333333
Recall: 0.2962962962962963
F1 score: 0.38095238095238093
```

### 1.3: Using Diabetes, High Blood Pressure and anaemia, platelets, smoking as predictor variables


```
In [26]: # print(data.head())
selected_predictor_variables = ['anaemia',
                               'diabetes',
                               'platelets', 'high_blood_pressure', 'smoking']
```

```
In [27]: classifier = DecisionTreeClassifier(criterion='entropy', random_state=0)
classifier.fit(x_train_scaled_df[selected_predictor_variables], y_train)
```

```
Out[27]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='entropy',
                                max_depth=None, max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort='deprecated',
                                random_state=0, splitter='best')
```

```
In [28]: y_pred = classifier.predict(x_val_scaled_df[selected_predictor_variables])
```

```
In [29]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [30]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy: 0.5066666666666667
Error: 0.4933333333333333
Precision: 0.2916666666666667
Recall: 0.25925925925925924
F1 score: 0.27450980392156865
```



#### 1.4: Using creatinine\_phosphokinase, serum\_creatinine, serum\_sodium and ejection\_fraction as predictor variables


```
In [31]: selected_predictor_variables = ['creatinine_phosphokinase',  
                                         'serum_creatinine',  
                                         'serum_sodium',  
                                         'ejection_fraction']
```

```
In [32]: classifier = DecisionTreeClassifier(criterion='entropy', random_state=0)  
classifier.fit(x_train_scaled_df[selected_predictor_variables], y_train)
```

```
Out[32]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='entropy',  
                                max_depth=None, max_features=None, max_leaf_nodes=None,  
                                min_impurity_decrease=0.0, min_impurity_split=None,  
                                min_samples_leaf=1, min_samples_split=2,  
                                min_weight_fraction_leaf=0.0, presort='deprecated',  
                                random_state=0, splitter='best')
```

```
In [33]: y_pred = classifier.predict(x_val_scaled_df[selected_predictor_variables])
```

```
In [34]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [35]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy: 0.6133333333333333
Error: 0.3866666666666667
Precision: 0.4642857142857143
Recall: 0.48148148148148145
F1 score: 0.4727272727272727
```

### 1.5: Using serum\_creatinine and ejection\_fraction as predictor variables


```
In [36]: selected_predictor_variables = ['serum_creatinine',  
                                         'ejection_fraction']
```

```
In [37]: classifier = DecisionTreeClassifier(criterion='entropy', random_state=0)  
classifier.fit(x_train_scaled_df[selected_predictor_variables], y_train)
```

```
Out[37]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='entropy',  
                                max_depth=None, max_features=None, max_leaf_nodes=None,  
                                min_impurity_decrease=0.0, min_impurity_split=None,  
                                min_samples_leaf=1, min_samples_split=2,  
                                min_weight_fraction_leaf=0.0, presort='deprecated',  
                                random_state=0, splitter='best')
```

```
In [38]: y_pred = classifier.predict(x_val_scaled_df[selected_predictor_variables])
```

```
In [39]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [40]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy: 0.72
Error: 0.28
Precision: 0.6363636363636364
Recall: 0.5185185185185185
F1 score: 0.5714285714285714
```

## 2) K Nearest Neighbors

### 2.1: Using all predictor variables


```
In [41]: classifier = KNeighborsClassifier(n_neighbors=3)
classifier.fit(x_train_scaled, y_train)
```

/home/ashwin/anaconda3/lib/python3.7/site-packages/ipykernel\_launcher.py:2: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
Out[41]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=None, n_neighbors=3, p=2,
                             weights='uniform')
```

```
In [42]: y_pred = classifier.predict(x_val_scaled)
```

```
In [43]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [44]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy:  0.6933333333333334
Error:  0.30666666666666664
Precision:  0.6666666666666666
Recall:  0.2962962962962963
F1 score:  0.41025641025641024
```

## 2.2: Using Age and Gender Variables

```
In [45]: selected_predictor_variables = ['age',  
                                         'sex']
```


```
In [46]: classifier = KNeighborsClassifier(n_neighbors=3)  
classifier.fit(x_train_scaled_df[selected_predictor_variables], y_train)
```

/home/ashwin/anaconda3/lib/python3.7/site-packages/ipykernel\_launcher.py:2: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
Out[46]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',  
                             metric_params=None, n_jobs=None, n_neighbors=3, p=2,  
                             weights='uniform')
```

```
In [47]: y_pred = classifier.predict(x_val_scaled_df[selected_predictor_variables])
```

```
In [48]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [49]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy: 0.5866666666666667
Error: 0.41333333333333333
Precision: 0.4166666666666667
Recall: 0.37037037037037035
F1 score: 0.39215686274509803
```



### 2.3: Using Diabetes, High Blood Pressure and anaemia, platelets, smoking as predictor variables

```
In [50]: selected_predictor_variables = ['anaemia',  
                                         'diabetes',  
                                         'platelets', 'high_blood_pressure', 'smoking']
```


```
In [51]: classifier = KNeighborsClassifier(n_neighbors=3)  
classifier.fit(x_train_scaled_df[selected_predictor_variables], y_train)
```

/home/ashwin/anaconda3/lib/python3.7/site-packages/ipykernel\_launcher.py:2: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
Out[51]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',  
                               metric_params=None, n_jobs=None, n_neighbors=3, p=2,  
                               weights='uniform')
```

```
In [52]: y_pred = classifier.predict(x_val_scaled_df[selected_predictor_variables])
```

```
In [53]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [54]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy: 0.5466666666666666
Error: 0.45333333333333337
Precision: 0.3333333333333333
Recall: 0.25925925925925924
F1 score: 0.2916666666666667
```

## 2.4: Using creatinine\_phosphokinase, serum\_creatinine, serum\_sodium and ejection\_fraction as predictor variables

```
In [55]: selected_predictor_variables = ['creatinine_phosphokinase',  
                                         'serum_creatinine',  
                                         'serum_sodium',  
                                         'ejection_fraction']
```


```
In [56]: classifier = KNeighborsClassifier(n_neighbors=3)  
classifier.fit(x_train_scaled_df[selected_predictor_variables], y_train)
```

/home/ashwin/anaconda3/lib/python3.7/site-packages/ipykernel\_launcher.py:2: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
Out[56]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',  
                               metric_params=None, n_jobs=None, n_neighbors=3, p=2,  
                               weights='uniform')
```

```
In [57]: y_pred = classifier.predict(x_val_scaled_df[selected_predictor_variables])
```

```
In [58]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [59]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy: 0.6666666666666666
Error: 0.33333333333333337
Precision: 0.5555555555555556
Recall: 0.37037037037037035
F1 score: 0.44444444444444445
```

## 2.5: Using serum\_creatinine and ejection\_fraction as predictor variables

```
In [60]: selected_predictor_variables = ['serum_creatinine',  
                                         'ejection_fraction']
```


```
In [61]: classifier_k_nearest = KNeighborsClassifier(n_neighbors=3)  
classifier_k_nearest.fit(x_train_scaled_df[selected_predictor_variables], y_train)
```

/home/ashwin/anaconda3/lib/python3.7/site-packages/ipykernel\_launcher.py:2: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
Out[61]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',  
                               metric_params=None, n_jobs=None, n_neighbors=3, p=2,  
                               weights='uniform')
```

```
In [62]: y_pred = classifier_k_nearest.predict(x_val_scaled_df[selected_predictor_variables])
```

```
In [63]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [64]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy: 0.64
Error: 0.36
Precision: 0.5
Recall: 0.4444444444444444
F1 score: 0.47058823529411764
```

## 3) Naive Bayes

### 3.1: Using all predictor variables


```
In [65]: classifier = GaussianNB()  
classifier.fit(x_train_scaled, y_train)
```

```
/home/ashwin/anaconda3/lib/python3.7/site-packages/sklearn/naive_bayes.py:206: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().  
  y = column_or_1d(y, warn=True)
```

```
Out[65]: GaussianNB(priors=None, var_smoothing=1e-09)
```

```
In [66]: y_pred = classifier.predict(x_val_scaled)
```

```
In [67]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [68]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy: 0.6666666666666666
Error: 0.33333333333333337
Precision: 0.5714285714285714
Recall: 0.2962962962962963
F1 score: 0.3902439024390244
```



### 3.2: Using Age and Gender Variables

```
In [69]: selected_predictor_variables = ['age',  
                                         'sex']
```

```
In [70]: classifier = GaussianNB()  
classifier.fit(x_train_scaled_df[selected_predictor_variables], y_train)
```


/home/ashwin/anaconda3/lib/python3.7/site-packages/sklearn/naive\_bayes.py:206: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

```
Out[70]: GaussianNB(priors=None, var_smoothing=1e-09)
```

```
In [71]: y_pred = classifier.predict(x_val_scaled_df[selected_predictor_variables])
```

```
In [72]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [73]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy: 0.6533333333333333
Error: 0.3466666666666667
Precision: 0.5714285714285714
Recall: 0.14814814814814814
F1 score: 0.23529411764705882
```

### 3.3: Using Diabetes, High Blood Pressure and anaemia, platelets, smoking as predictor variables

```
In [74]: selected_predictor_variables = ['anaemia',  
                                         'diabetes',  
                                         'platelets', 'high_blood_pressure', 'smoking']
```

```
In [75]: classifier = GaussianNB()  
classifier.fit(x_train_scaled_df[selected_predictor_variables], y_train)
```


/home/ashwin/anaconda3/lib/python3.7/site-packages/sklearn/naive\_bayes.py:206: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

```
Out[75]: GaussianNB(priors=None, var_smoothing=1e-09)
```

```
In [76]: y_pred = classifier.predict(x_val_scaled_df[selected_predictor_variables])
```

```
In [77]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [78]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy:  0.64
Error:  0.36
Precision:  0.0
Recall:  0.0
F1 score:  0.0
```

```
/home/ashwin/anaconda3/lib/python3.7/site-packages/sklearn/metrics/_classification.py:1272: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

### 3.4: Using creatinine\_phosphokinase, serum\_creatinine, serum\_sodium and ejection\_fraction as predictor variables

```
In [79]: selected_predictor_variables = ['creatinine_phosphokinase',
                                         'serum_creatinine',
                                         'serum_sodium',
                                         'ejection_fraction']
```

```
In [80]: classifier = GaussianNB()
classifier.fit(x_train_scaled_df[selected_predictor_variables], y_train)
```


```
/home/ashwin/anaconda3/lib/python3.7/site-packages/sklearn/naive_bayes.py:206: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
```

```
y = column_or_1d(y, warn=True)
```

```
Out[80]: GaussianNB(priors=None, var_smoothing=1e-09)
```

```
In [81]: y_pred = classifier.predict(x_val_scaled_df[selected_predictor_variables])
```

```
In [82]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [83]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy:  0.64
Error:  0.36
Precision:  0.5
Recall:  0.18518518518518517
F1 score:  0.2702702702702703
```

### 3.5: Using serum\_creatinine and ejection\_fraction as predictor variables

```
In [84]: selected_predictor_variables = ['serum_creatinine',
                                         'ejection_fraction']
```


```
In [85]: classifier = GaussianNB()
classifier.fit(x_train_scaled_df[selected_predictor_variables], y_train)
```

```
/home/ashwin/anaconda3/lib/python3.7/site-packages/sklearn/naive_bayes.py:206: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
```

```
Out[85]: GaussianNB(priors=None, var_smoothing=1e-09)
```

```
In [86]: y_pred = classifier.predict(x_val_scaled_df[selected_predictor_variables])
```

```
In [87]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [88]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy: 0.6666666666666666
Error: 0.33333333333333337
Precision: 0.6666666666666666
Recall: 0.14814814814814814
F1 score: 0.24242424242424243
```



## 4) Support Vector Machine

### 4.1: Using all predictor variables


```
In [89]: classifier = SVC(kernel='rbf')  
classifier.fit(x_train_scaled, y_train)
```

```
/home/ashwin/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:760: DataConversionWarning:  
A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ),  
for example using ravel().  
y = column_or_1d(y, warn=True)
```

```
Out[89]: SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,  
decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',  
max_iter=-1, probability=False, random_state=None, shrinking=True,  
tol=0.001, verbose=False)
```

```
In [90]: y_pred = classifier.predict(x_val_scaled)
```

```
In [91]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [92]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy:  0.8133333333333334
Error:  0.18666666666666665
Precision:  0.8421052631578947
Recall:  0.5925925925925926
F1 score:  0.6956521739130435
```

## 4.2: Using Age and Gender Variables

```
In [93]: selected_predictor_variables = ['age',  
                                         'sex']
```

```
In [94]: classifier = SVC(kernel='rbf')  
classifier.fit(x_train_scaled_df[selected_predictor_variables], y_train)
```


/home/ashwin/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:760: DataConversionWarning:  
A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ),  
for example using ravel().

```
y = column_or_1d(y, warn=True)
```

```
Out[94]: SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,  
             decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',  
             max_iter=-1, probability=False, random_state=None, shrinking=True,  
             tol=0.001, verbose=False)
```

```
In [95]: y_pred = classifier.predict(x_val_scaled_df[selected_predictor_variables])
```

```
In [96]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [97]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy: 0.6533333333333333
Error: 0.3466666666666667
Precision: 0.5714285714285714
Recall: 0.14814814814814814
F1 score: 0.23529411764705882
```

#### 4.3: Using Diabetes, High Blood Pressure and anaemia, platelets, smoking as predictor variables

```
In [98]: selected_predictor_variables = ['anaemia',  
                                       'diabetes',  
                                       'platelets', 'high_blood_pressure', 'smoking']
```

```
In [99]: classifier = SVC(kernel='rbf')  
classifier.fit(x_train_scaled_df[selected_predictor_variables], y_train)
```


/home/ashwin/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:760: DataConversionWarning:  
A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ),  
for example using ravel().

```
y = column_or_1d(y, warn=True)
```

```
Out[99]: SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,  
            decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',  
            max_iter=-1, probability=False, random_state=None, shrinking=True,  
            tol=0.001, verbose=False)
```

```
In [100]: y_pred = classifier.predict(x_val_scaled_df[selected_predictor_variables])
```

```
In [101]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```





```
In [104]: classifier = SVC(kernel='rbf')
classifier.fit(x_train_scaled_df[selected_predictor_variables], y_train)
```


/home/ashwin/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:760: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

```
Out[104]: SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',
max_iter=-1, probability=False, random_state=None, shrinking=True,
tol=0.001, verbose=False)
```

```
In [105]: y_pred = classifier.predict(x_val_scaled_df[selected_predictor_variables])
```

```
In [106]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```





```
In [107]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy:  0.7333333333333333
Error:  0.26666666666666667
Precision:  0.6842105263157895
Recall:  0.48148148148148145
F1 score:  0.5652173913043478
```

#### 4.5: Using serum\_creatinine and ejection\_fraction as predictor variables

```
In [108]: selected_predictor_variables = ['serum_creatinine',
                                         'ejection_fraction']
```


```
In [109]: classifier = SVC(kernel='rbf')
classifier.fit(x_train_scaled_df[selected_predictor_variables], y_train)
```

```
/home/ashwin/anaconda3/lib/python3.7/site-packages/sklearn/utils/validation.py:760: DataConversionWarning:
A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ),
for example using ravel().
  y = column_or_1d(y, warn=True)
```

```
Out[109]: SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
              decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',
              max_iter=-1, probability=False, random_state=None, shrinking=True,
              tol=0.001, verbose=False)
```

```
In [110]: y_pred = classifier.predict(x_val_scaled_df[selected_predictor_variables])
```

```
In [111]: conf_matrix = metrics.confusion_matrix(y_val, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = ".3f", square = True, cmap = plt.cm.Blues)
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.title('Confusion matrix')
plt.tight_layout()
```



```
In [112]: accuracy = metrics.accuracy_score(y_val, y_pred)
print("Accuracy: ", accuracy)
error = 1 - accuracy
print("Error: ", error)
precision = metrics.precision_score(y_val, y_pred)
print("Precision: ", precision)
recall = metrics.recall_score(y_val, y_pred)
print("Recall: ", recall)
F1_score = metrics.f1_score(y_val, y_pred)
print("F1 score: ", F1_score)
```

```
Accuracy:  0.7466666666666667
Error:  0.25333333333333333
Precision:  0.7
Recall:  0.5185185185185185
F1 score:  0.5957446808510639
```

# Classification: Model Performances

- Decision Tree

Variables	Accuracy	Precision	Recall	F1 Score
All	0.81	0.72	0.78	0.75
Age & Gender	0.65	0.53	0.30	0.38
Diabetes, High BP, Platelets, Anaemia & Smoking	0.51	0.29	0.26	0.27
creatinine_phosphokinase, serum_creatinine,serum_sodium and ejection_fraction	0.61	0.46	0.48	0.47
serum_sodium and ejection_fraction	0.72	0.64	0.52	0.57

- K Nearest Neighbors

Variables	Accuracy	Precision	Recall	F1 Score
All	0.69	0.67	0.30	0.41
Age & Gender	0.59	0.42	0.37	0.39
Diabetes, High BP, Platelets, Anaemia & Smoking	0.55	0.33	0.26	0.29
creatinine_phosphokinase, serum_creatinine,serum_sodium and ejection_fraction	0.67	0.56	0.37	0.44
serum_sodium and ejection_fraction	0.64	0.50	0.44	0.47

- Naive Bayes

Variables	Accuracy	Precision	Recall	F1 Score
All	0.67	0.57	0.30	0.39
Age & Gender	0.65	0.57	0.15	0.24
Diabetes, High BP, Platelets, Anaemia & Smoking	0.64	0.00	0.00	0.00
creatinine_phosphokinase, serum_creatinine,serum_sodium and ejection_fraction	0.64	0.50	0.19	0.27
serum_sodium and ejection_fraction	0.67	0.67	0.15	0.24

- Support Vector Machines

Variables	Accuracy	Precision	Recall	F1 Score
All	0.81	0.84	0.59	0.70
Age & Gender	0.65	0.57	0.15	0.24
Diabetes, High BP, Platelets, Anaemia & Smoking	0.64	0.00	0.00	0.00
creatinine_phosphokinase, serum_creatinine,serum_sodium and ejection_fraction	0.73	0.68	0.48	0.57
serum_sodium and ejection_fraction	0.75	0.70	0.52	0.60

### Best performing classification model

Decision Tree using all predictor variables gives the best F1 Score

### The performance of the model

- Accuracy: 0.81
- Precision: 0.72
- Recall: 0.78
- F1 Score: 0.75

### Selection of the parameters of the model

- Upon trial of many different combinations we understand that few group of variables do not contribute to the best performance of a model. So we chose all predictor variables for the classification as they give the best F1 score and performance.

```
In [113]: # Best performing classifiers & variables\n",
# (DecisionTreeClassifier(criterion='entropy', random_state=0), ('age', 'anaemia', 'creatinine_phosphokinase', 'platelets', 'serum_creatinine', 'serum_sodium', 'smoking', 'time'))\n",
# (DecisionTreeClassifier(criterion='entropy', random_state=0), ('age', 'creatinine_phosphokinase', 'diabetes', 'platelets', 'serum_creatinine', 'serum_sodium', 'smoking', 'time'))\n",
# (DecisionTreeClassifier(criterion='entropy', random_state=0), ('creatinine_phosphokinase', 'diabetes', 'ejection_fraction', 'high_blood_pressure', 'platelets', 'serum_creatinine', 'serum_sodium', 'smoking', 'time'))\n",
# (DecisionTreeClassifier(criterion='entropy', random_state=0), ('ejection_fraction', 'serum_creatinine', 'serum_sodium', 'time'))\n",
# (DecisionTreeClassifier(criterion='entropy', random_state=0), ('ejection_fraction', 'high_blood_pressure', 'serum_creatinine', 'serum_sodium', 'time'))"
```

## 5 Clustering

```
In [114]: data = pd.read_csv('heart_failure_clinical_records_dataset.csv')
data.head()


x_train = data.iloc[:, :-1]
y_train = data.iloc[:, -1].to_numpy().flatten()

scaler = StandardScaler()
scaler.fit(x_train)
x_train_scaled = scaler.transform(x_train)
x_train_scaled_df = pd.DataFrame(x_train_scaled, index=x_train.index, columns=x_train.columns)
```

Following are the two true clusters based on "Serum Creatinine" and "Ejection Fraction"

```
In [115]: # clustering = linkage(all_x, method = "single", metric = "euclidean")
# clusters = fcluster(clustering, 2, criterion = 'maxclust')
col1 = 'serum_creatinine'
col2 = 'ejection_fraction'
ax = sns.scatterplot(data=data, x=col1, y=col2, hue='DEATH_EVENT')
ax.set(title = '{} vs {}'.format(col1, col2), xlabel = col1, ylabel = col2)
```

```
Out[115]: [Text(0, 0.5, 'ejection_fraction'),
Text(0.5, 0, 'serum_creatinine'),
Text(0.5, 1.0, 'serum_creatinine vs ejection_fraction')]
```



## 5.1 Hierarchical Clustering - Single Linkage

### 5.1.1 All variables

```
In [116]: all_x = x_train_scaled_df.to_numpy()
```


```
In [117]: clustering = linkage(all_x, method="single", metric = "euclidean")
clusters = fcluster(clustering, 2, criterion = 'maxclust')
```

```
In [118]: adjusted_rand_index = metrics.adjusted_rand_score(y_train, clusters)
silhouette_coefficient = metrics.silhouette_score(all_x, clusters, metric = "euclidean")
print([adjusted_rand_index, silhouette_coefficient])
```

```
[0.007441167976633384, 0.41639364203172197]
```

```
In [119]: data['clusters'] = clusters - 1
ax = sns.scatterplot(data=data, x=col1, y=col2, hue='clusters')
ax.set(title = '{} vs {}'.format(col1, col2), xlabel = col1, ylabel = col2)
```

```
Out[119]: [Text(0, 0.5, 'ejection_fraction'),
Text(0.5, 0, 'serum_creatinine'),
Text(0.5, 1.0, 'serum_creatinine vs ejection_fraction')]
```



### 5.1.2: Using Age and Gender Variables




```
In [120]: selected_predictor_variables = ['age',  
                                         'sex']  
x = x_train_scaled_df.loc[:, selected_predictor_variables].to_numpy()
```

```
In [121]: clustering = linkage(x, method="single", metric = "euclidean")  
clusters = fcluster(clustering, 2, criterion = 'maxclust')
```

```
In [122]: adjusted_rand_index = metrics.adjusted_rand_score(y_train, clusters)  
silhouette_coefficient = metrics.silhouette_score(x, clusters, metric = "euclidean")  
print([adjusted_rand_index, silhouette_coefficient])  
  
[-0.0018582102712868956, 0.551937322742283]
```

```
In [123]: data['clusters'] = clusters - 1  
ax = sns.scatterplot(data=data, x=col1, y=col2, hue='clusters')  
ax.set(title = '{} vs {}'.format(col1, col2), xlabel = col1, ylabel = col2)
```

```
Out[123]: [Text(0, 0.5, 'ejection_fraction'),  
           Text(0.5, 0, 'serum_creatinine'),  
           Text(0.5, 1.0, 'serum_creatinine vs ejection_fraction')]
```



### 5.1.3: Using Health Problem Variables


```
In [124]: selected_predictor_variables = ['anaemia',  
                                         'diabetes',  
                                         'platelets', 'high_blood_pressure', 'smoking']  
x = x_train_scaled_df.loc[:, selected_predictor_variables].to_numpy()
```

```
In [125]: clustering = linkage(x, method="single", metric = "euclidean")  
clusters = fcluster(clustering, 2, criterion = 'maxclust')
```

```
In [126]: adjusted_rand_index = metrics.adjusted_rand_score(y_train, clusters)  
silhouette_coefficient = metrics.silhouette_score(x, clusters, metric = "euclidean")  
print([adjusted_rand_index, silhouette_coefficient])  
[-0.0035189759889497152, 0.5427711036936387]
```

```
In [127]: data['clusters'] = clusters - 1
ax = sns.scatterplot(data=data, x=col1, y=col2, hue='clusters')
ax.set(title = '{} vs {}'.format(col1, col2), xlabel = col1, ylabel = col2)
```

```
Out[127]: [Text(0, 0.5, 'ejection_fraction'),
Text(0.5, 0, 'serum_creatinine'),
Text(0.5, 1.0, 'serum_creatinine vs ejection_fraction')]
```



#### 5.1.4: Using Body Parameter Variables

```
In [128]: selected_predictor_variables = ['creatinine_phosphokinase',
                                          'serum_creatinine',
                                          'serum_sodium',
                                          'ejection_fraction']
x = x_train_scaled_df.loc[:, selected_predictor_variables].to_numpy()
```


```
In [129]: clustering = linkage(x, method="single", metric = "euclidean")
clusters = fcluster(clustering, 2, criterion = 'maxclust')
```

```
In [130]: adjusted_rand_index = metrics.adjusted_rand_score(y_train, clusters)
silhouette_coefficient = metrics.silhouette_score(x, clusters, metric = "euclidean")
print([adjusted_rand_index, silhouette_coefficient])
```

```
[0.007441167976633384, 0.6662613003976123]
```

```
In [131]: data['clusters'] = clusters - 1
ax = sns.scatterplot(data=data, x=col1, y=col2, hue='clusters')
ax.set(title = '{} vs {}'.format(col1, col2), xlabel = col1, ylabel = col2)
```

```
Out[131]: [Text(0, 0.5, 'ejection_fraction'),
Text(0.5, 0, 'serum_creatinine'),
Text(0.5, 1.0, 'serum_creatinine vs ejection_fraction')]
```



### 5.1.5: Using Anaemia

```
In [132]: selected_predictor_variables = ['anaemia']
x = x_train_scaled_df.loc[:, selected_predictor_variables].to_numpy()
```


```
In [133]: clustering = linkage(x, method="single", metric = "euclidean")
clusters = fcluster(clustering, 2, criterion = 'maxclust')
```

```
In [134]: adjusted_rand_index = metrics.adjusted_rand_score(y_train, clusters)
silhouette_coefficient = metrics.silhouette_score(x, clusters, metric = "euclidean")
print([adjusted_rand_index, silhouette_coefficient])

[0.006938789172657143, 1.0]
```

```
In [135]: data['clusters'] = clusters - 1
ax = sns.scatterplot(data=data, x=col1, y=col2, hue='clusters')
ax.set(title = '{} vs {}'.format(col1, col2), xlabel = col1, ylabel = col2)
```

```
Out[135]: [Text(0, 0.5, 'ejection_fraction'),
Text(0.5, 0, 'serum_creatinine'),
Text(0.5, 1.0, 'serum_creatinine vs ejection_fraction')]
```



## 2) Hierarchical Clustering - Complete Linkage

### 5.2.1 All variables

```
In [136]: all_x = x_train_scaled_df.to_numpy()
```


```
In [137]: clustering = linkage(all_x, method="complete", metric = "euclidean")
clusters = fcluster(clustering, 2, criterion = 'maxclust')
```

```
In [138]: adjusted_rand_index = metrics.adjusted_rand_score(y_train, clusters)
silhouette_coefficient = metrics.silhouette_score(all_x, clusters, metric = "euclidean")
print([adjusted_rand_index, silhouette_coefficient])

[0.022507236951381675, 0.46861144710828245]
```

```
In [139]: data['clusters'] = clusters - 1
ax = sns.scatterplot(data=data, x=col1, y=col2, hue='clusters')
ax.set(title = '{} vs {}'.format(col1, col2), xlabel = col1, ylabel = col2)
```

```
Out[139]: [Text(0, 0.5, 'ejection_fraction'),
Text(0.5, 0, 'serum_creatinine'),
Text(0.5, 1.0, 'serum_creatinine vs ejection_fraction')]
```



### 5.2.2: Using Age and Gender Variables


```
In [140]: selected_predictor_variables = ['age',  
                                         'sex']  
x = x_train_scaled_df.loc[:, selected_predictor_variables].to_numpy()
```

```
In [141]: clustering = linkage(x, method="complete", metric = "euclidean")  
clusters = fcluster(clustering, 2, criterion = 'maxclust')
```

```
In [142]: adjusted_rand_index = metrics.adjusted_rand_score(y_train, clusters)  
silhouette_coefficient = metrics.silhouette_score(x, clusters, metric = "euclidean")  
print([adjusted_rand_index, silhouette_coefficient])  
  
[0.11142775349398941, 0.3419871555832537]
```

```
In [143]: data['clusters'] = clusters - 1
ax = sns.scatterplot(data=data, x=col1, y=col2, hue='clusters')
ax.set(title = '{} vs {}'.format(col1, col2), xlabel = col1, ylabel = col2)
```

```
Out[143]: [Text(0, 0.5, 'ejection_fraction'),
Text(0.5, 0, 'serum_creatinine'),
Text(0.5, 1.0, 'serum_creatinine vs ejection_fraction')]
```



### 5.2.3: Using Health Problem Variables

```
In [144]: selected_predictor_variables = ['anaemia',
                                          'diabetes',
                                          'platelets', 'high_blood_pressure', 'smoking']
x = x_train_scaled_df.loc[:, selected_predictor_variables].to_numpy()
```

```
In [145]: clustering = linkage(x, method="complete", metric = "euclidean")
clusters = fcluster(clustering, 2, criterion = 'maxclust')
```




```
In [146]: adjusted_rand_index = metrics.adjusted_rand_score(y_train, clusters)
silhouette_coefficient = metrics.silhouette_score(x, clusters, metric = "euclidean")
print([adjusted_rand_index, silhouette_coefficient])
```

```
[-0.007387031337951726, 0.11014816465932352]
```

```
In [147]: data['clusters'] = clusters - 1
ax = sns.scatterplot(data=data, x=col1, y=col2, hue='clusters')
ax.set(title = '{} vs {}'.format(col1, col2), xlabel = col1, ylabel = col2)
```

```
Out[147]: [Text(0, 0.5, 'ejection_fraction'),
Text(0.5, 0, 'serum_creatinine'),
Text(0.5, 1.0, 'serum_creatinine vs ejection_fraction')]
```



## 5.2.4: Using Body Parameter Variables


```
In [148]: selected_predictor_variables = ['creatinine_phosphokinase',  
                                         'serum_creatinine',  
                                         'serum_sodium',  
                                         'ejection_fraction']  
x = x_train_scaled_df.loc[:, selected_predictor_variables].to_numpy()
```

```
In [149]: clustering = linkage(x, method="complete", metric = "euclidean")  
clusters = fcluster(clustering, 2, criterion = 'maxclust')
```

```
In [150]: adjusted_rand_index = metrics.adjusted_rand_score(y_train, clusters)  
silhouette_coefficient = metrics.silhouette_score(x, clusters, metric = "euclidean")  
print([adjusted_rand_index, silhouette_coefficient])  
  
[0.02612613619798278, 0.6567989004211781]
```

```
In [151]: data['clusters'] = clusters - 1
ax = sns.scatterplot(data=data, x=col1, y=col2, hue='clusters')
ax.set(title = '{} vs {}'.format(col1, col2), xlabel = col1, ylabel = col2)
```

```
Out[151]: [Text(0, 0.5, 'ejection_fraction'),
Text(0.5, 0, 'serum_creatinine'),
Text(0.5, 1.0, 'serum_creatinine vs ejection_fraction')]
```



## 5.2.5: Using Anaemia

```
In [152]: selected_predictor_variables = ['anaemia']
x = x_train_scaled_df.loc[:, selected_predictor_variables].to_numpy()
```


```
In [153]: clustering = linkage(x, method="complete", metric = "euclidean")
clusters = fcluster(clustering, 2, criterion = 'maxclust')
```

```
In [154]: adjusted_rand_index = metrics.adjusted_rand_score(y_train, clusters)
          silhouette_coefficient = metrics.silhouette_score(x, clusters, metric = "euclidean")
          print([adjusted_rand_index, silhouette_coefficient])

[0.006938789172657143, 1.0]
```

```
In [155]: data['clusters'] = clusters - 1
          ax = sns.scatterplot(data=data, x=col1, y=col2, hue='clusters')
          ax.set(title = '{} vs {}'.format(col1, col2), xlabel = col1, ylabel = col2)
```

```
Out[155]: [Text(0, 0.5, 'ejection_fraction'),
           Text(0.5, 0, 'serum_creatinine'),
           Text(0.5, 1.0, 'serum_creatinine vs ejection_fraction')]
```



### 3) KMeans Clustering

#### 5.3.1: All variables

```
In [156]: all_x = x_train_scaled_df.to_numpy()
```


```
In [157]: clustering = KMeans(n_clusters = 2, init = 'random', n_init = 10, random_state=0).fit(all_x)
clusters = clustering.labels_
```

```
In [158]: adjusted_rand_index = metrics.adjusted_rand_score(y_train, clusters)
silhouette_coefficient = metrics.silhouette_score(all_x, clusters, metric = "euclidean")
print([adjusted_rand_index, silhouette_coefficient])
```

```
[-0.0036799581865726927, 0.1178512561512019]
```

```
In [159]: data['clusters'] = clusters
ax = sns.scatterplot(data=data, x=col1, y=col2, hue='clusters')
ax.set(title = '{} vs {}'.format(col1, col2), xlabel = col1, ylabel = col2)
```

```
Out[159]: [Text(0, 0.5, 'ejection_fraction'),
Text(0.5, 0, 'serum_creatinine'),
Text(0.5, 1.0, 'serum_creatinine vs ejection_fraction')]
```



### 5.3.2 Using Age and Gender Variables


```
In [160]: selected_predictor_variables = ['age',  
                                         'sex']  
x = x_train_scaled_df.loc[:, selected_predictor_variables].to_numpy()
```

```
In [161]: clustering = KMeans(n_clusters = 2, init = 'random', n_init = 10, random_state=0).fit(x)  
clusters = clustering.labels_
```

```
In [162]: adjusted_rand_index = metrics.adjusted_rand_score(y_train, clusters)  
silhouette_coefficient = metrics.silhouette_score(x, clusters, metric = "euclidean")  
print([adjusted_rand_index, silhouette_coefficient])  
  
[-0.0018582102712868956, 0.551937322742283]
```

```
In [163]: data['clusters'] = clusters  
ax = sns.scatterplot(data=data, x=col1, y=col2, hue='clusters')  
ax.set(title = '{} vs {}'.format(col1, col2), xlabel = col1, ylabel = col2)
```

```
Out[163]: [Text(0, 0.5, 'ejection_fraction'),  
Text(0.5, 0, 'serum_creatinine'),  
Text(0.5, 1.0, 'serum_creatinine vs ejection_fraction')]
```



### 5.3.3 Using Health Problem Variables


```
In [164]: selected_predictor_variables = ['anaemia',  
                                         'diabetes',  
                                         'platelets', 'high_blood_pressure', 'smoking']  
x = x_train_scaled_df.loc[:, selected_predictor_variables].to_numpy()
```

```
In [165]: clustering = KMeans(n_clusters = 2, init = 'random', n_init = 10, random_state=0).fit(x)  
clusters = clustering.labels_
```

```
In [166]: adjusted_rand_index = metrics.adjusted_rand_score(y_train, clusters)  
silhouette_coefficient = metrics.silhouette_score(x, clusters, metric = "euclidean")  
print([adjusted_rand_index, silhouette_coefficient])  
[-0.005350592860670945, 0.24587001478483037]
```

```
In [167]: data['clusters'] = clusters
ax = sns.scatterplot(data=data, x=col1, y=col2, hue='clusters')
ax.set(title = '{} vs {}'.format(col1, col2), xlabel = col1, ylabel = col2)
```

```
Out[167]: [Text(0, 0.5, 'ejection_fraction'),
Text(0.5, 0, 'serum_creatinine'),
Text(0.5, 1.0, 'serum_creatinine vs ejection_fraction')]
```



### 5.3.4 Using Body Parameter Variables

```
In [168]: selected_predictor_variables = ['creatinine_phosphokinase',
                                          'serum_creatinine',
                                          'serum_sodium',
                                          'ejection_fraction']
x = x_train_scaled_df.loc[:, selected_predictor_variables].to_numpy()
```

```
In [169]: clustering = KMeans(n_clusters = 2, init = 'random', n_init = 10, random_state=0).fit(x)
clusters = clustering.labels_
```




```
In [170]: adjusted_rand_index = metrics.adjusted_rand_score(y_train, clusters)
silhouette_coefficient = metrics.silhouette_score(x, clusters, metric = "euclidean")
print([adjusted_rand_index, silhouette_coefficient])
```

```
[0.15419143591523968, 0.24027667508494274]
```

```
In [171]: data['clusters'] = clusters
ax = sns.scatterplot(data=data, x=col1, y=col2, hue='clusters')
ax.set(title = '{} vs {}'.format(col1, col2), xlabel = col1, ylabel = col2)
```

```
Out[171]: [Text(0, 0.5, 'ejection_fraction'),
Text(0.5, 0, 'serum_creatinine'),
Text(0.5, 1.0, 'serum_creatinine vs ejection_fraction')]
```



### 5.3.5 Using Anaemia

```
In [172]: selected_predictor_variables = ['anaemia']
x = x_train_scaled_df.loc[:, selected_predictor_variables].to_numpy()
```


```
In [173]: clustering = KMeans(n_clusters = 2, init = 'random', n_init = 10, random_state=0).fit(x)
clusters = clustering.labels_
```

```
In [174]: adjusted_rand_index = metrics.adjusted_rand_score(y_train, clusters)
silhouette_coefficient = metrics.silhouette_score(x, clusters, metric = "euclidean")
print([adjusted_rand_index, silhouette_coefficient])

[0.006938789172657143, 1.0]
```

```
In [175]: data['clusters'] = clusters
ax = sns.scatterplot(data=data, x=col1, y=col2, hue='clusters')
ax.set(title = '{} vs {}'.format(col1, col2), xlabel = col1, ylabel = col2)
```

```
Out[175]: [Text(0, 0.5, 'ejection_fraction'),
Text(0.5, 0, 'serum_creatinine'),
Text(0.5, 1.0, 'serum_creatinine vs ejection_fraction')]
```



## Classification: Model Performances

- Hierarchical Clustering - Single Linkage

Variables	Rand Index	Silhouette Coefficient
All	0.007	0.416
Age & Gender	-0.001	0.552
Health Problem	-0.003	0.543
Body Parameters	0.007	0.666
Anaemia	0.007	1.0

- Hierarchical Clustering - Complete Linkage

Variables	Rand Index	Silhouette Coefficient
All	0.022	0.469
Age & Gender	0.111	0.342
Health Problem	-0.007	0.11
Body Parameters	0.026	0.657
Anaemia	0.007	1.0

- KMeans Clustering

Variables	Rand Index	Silhouette Coefficient
All	-0.003	0.117
Age & Gender	-0.001	0.552
Health Problem	-0.005	0.246
Body Parameters	0.154	0.240
Anaemia	0.007	1.0

In [ ]: